

# 7 Code Shape

## Chapter Contents

1. Assigning Storage Locations
2. Arithmetic Operators
3. Boolean and Relational Operators
4. Storing and Accessing Arrays
5. Character Strings
6. Structure References
7. Control-Flow Constructs
8. Procedure Calls

## Code Shape

Typically, there are **many different ways** how a compiler can map a source-language construct into a sequence of operations in the target instruction set of a given processor.

These variations use different operations and different approaches

- some of these implementations are faster than others
- some use less memory
- some use fewer registers
- some might consume less energy during execution

The concept of **code shape** encapsulates all of the decisions, large and small, that the compiler writer makes on how to represent the computation in both IR and assembly code.

Code shape has a **strong impact** both on the behavior of the compiled code and on the ability of the optimizer and back end to improve it.

## Code Shape

**Example** Consider the way that a C compiler might implement a `switch` statement that switched on a single-byte character value.

```
1 char ch = ...;
2 switch(ch) {
3     key1: ... break;
4     ...
5     key256: ... break;
6     default: ...;
7 }
```

### Alternatives

1. use a cascaded series of `if-then-else` statements to implement the `switch` statement
2. use tests that perform a binary search
3. construct a table of 256 labels and interpret the character by loading the corresponding table entry and jumping to it

Which implementation is best for a particular `switch` statement depends on many factors

- number of cases and their relative execution frequencies
- cost structure for branching on the processor

## Assigning Storage Locations

The compiler must **assign a storage location** to each value produced by the code, taking the following factors into account.

- the value's type, its size, its visibility, and its lifetime
- runtime layout of memory
- source-language constraints on the layout of data areas and data structures
- target-processor constraints on placement or use of data

We can distinguish **two types** of values

- **named value**: lifetime is defined by source-language rules and actual use in the code
- **unnamed values**: must be handled consistently with the meaning of the program, but the compiler has great leeway in determining where these values reside and how long to retain them.

## Assigning Storage Locations

For each value, the compiler must also decide whether to keep it in a **register** or to keep it in **memory**. The compiler's **memory model** guides its choice of locations for values.

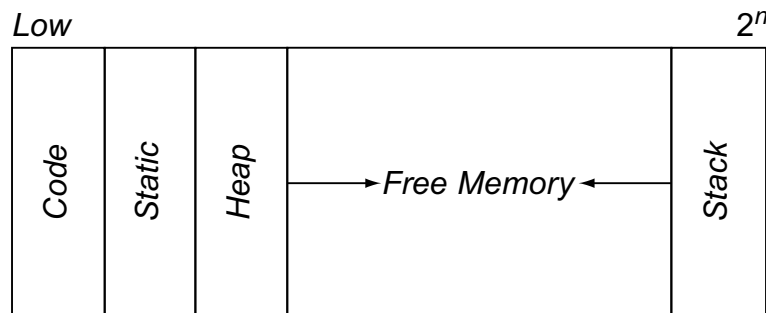
### Memory-to-Memory Model

- compiler assumes that all values reside in memory
- values are loaded into registers as needed and stored to memory after each definition
- IR typically uses **physical register** names

### Register-to-Register Model

- compiler assumes that it has enough registers to express the computation
- use a distinct **virtual register** for each value that can legally reside in a register
- store a virtual register's value to memory only when absolutely necessary

## Placing Runtime Data Structures

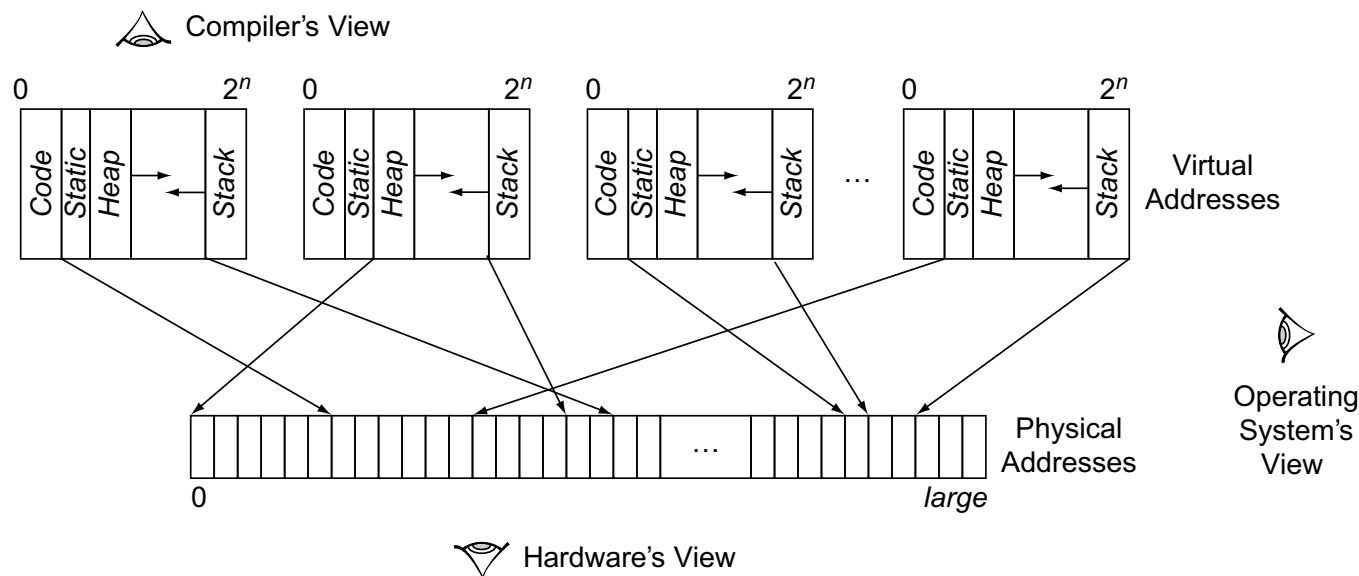


### Typical layout for the address space used by a single compiled program

- **code** sits at the bottom of the address space
- **static area** holds both static and global data areas, along with any fixed-size data created by the compiler, e.g., jump tables, debug information
- **heap** contains dynamically allocated data structures
- runtime **stack** is used to stack-allocate activation records, if possible

## Placing Runtime Data Structures

The operating system maps multiple **logical address spaces** into the single **physical address space** supported by the processor.



## Layout for Data Areas

For convenience, the compiler groups together the storage for values with the **same lifetimes** and **visibility**. It creates distinct **data areas** for them.

### Typical placement rules for Algol-like languages

- if  $x$  is declared locally in procedure  $p$ 
  - if its value is not preserved across distinct invocations of  $p$   
→ assign it to procedure-local storage
  - if its value is preserved across invocations of  $p$   
→ assign it to procedure-local static storage
- if  $x$  is declared as globally visible  
→ assign it to global storage
- if  $x$  is allocated under program control  
→ assign it to the runtime heap



## Assigning Offsets

In the case of local, static, and global data areas, the compiler must assign each name an **offset** inside the data area. Target ISAs constrain the placement of data items in memory.

### Typical alignment rules

- 32-bit integers and floating-point numbers begin on word (32-bit) boundaries
- 64-bit integer and floating-point data begin on doubleword (64-bit) boundaries
- string data begin on halfword (16-bit) boundaries

To **minimize wasted space**, the compiler should order the variables into groups, from those with the most restrictive alignment rules to those with the least.

**Note** Doubleword alignment is more restrictive than word alignment.

## Keeping Values in Registers

The register-to-register memory model has **three** principal advantages

- it is simple
- it can improve the results of analysis and optimization
- it enhances portability by postponing processor-specific constraints until optimization

However, only **unambiguous values** can be kept in registers. **Ambiguous values** cannot be kept in a register across either a definition or a use of another ambiguous value.

### Ambiguous and Unambiguous Values

Any value that can be accessed by multiple names is **ambiguous**. In contrast, a value that can be accessed with just one name is **unambiguous**.

## Keeping Values in Registers

### Ambiguity arises in several ways

- values stored in pointer-based variables are often ambiguous
- call-by-reference formal parameters and name scoping rules can make the formal parameters ambiguous
- array-element values can be ambiguous values since two references could refer to the same location

$$\begin{aligned} a &\leftarrow m + n \\ b &\leftarrow 13 \\ c &\leftarrow a + b \end{aligned}$$

With careful analysis, the compiler can **disambiguate** some of these cases.

- if  $a$  and  $b$  refer to the same location,  $c$  gets the value 26, otherwise it gets  $m + n + 13$
- to keep  $a$  in a register across assignment to another ambiguous value, the compiler needs to prove that the sets of locations to which  $a$  and  $b$  can refer are **disjoint**

Since pairwise ambiguity analysis is expensive, compilers typically relegate ambiguous values to memory, with a load before each use and a store after each definition.

## Arithmetic Operators

Modern processors provide a full complement of operations to evaluate expressions

- arithmetic operators, e.g., add, sub, imul, and idiv
- shift and rotate operators, e.g., shl, shr, rol, and ror
- boolean operators, e.g., and, or, xor, and not

To generate code for a **trivial expression**, such as  $a + b$ , the compiler emits code to load the values of  $a$  and  $b$  into registers, followed by an instruction to perform the addition.

```
loadI  @a          ⇒ r1
loadA0 rarp, r1 ⇒ ra
loadI  @b          ⇒ r2
loadA0 rarp, r2 ⇒ rb
add     ra, rb   ⇒ rt
```

If the value of  $a$  or  $b$  is already in a register, the compiler can avoid the load instructions and simply use that register in place of  $r_a$  or  $r_b$ , respectively.

## Arithmetic Operators

If the expression is represented in a tree-like IR, this process fits into a **postorder** tree walk.

### base

- returns the name of a register holding the base address for an identifier
- if needed, it emits code to get that address into a register

### offset

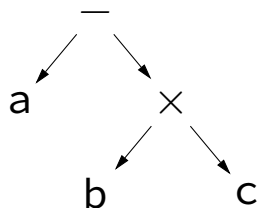
- returns the name of a register holding the identifier's offset
- offset is relative to the address returned by base

```

1 procedure expr(n)
2   if  $n \in \{+, -, \times, \div\}$  then
3      $t_1 \leftarrow \text{expr}(n.\text{left})$ 
4      $t_2 \leftarrow \text{expr}(n.\text{right})$ 
5      $r \leftarrow \text{NextRegister}()$ 
6     emit(n,  $t_1$ ,  $t_2$ , r)
7   else if  $n = \text{ident}$  then
8      $t_1 \leftarrow \text{base}(n)$ 
9      $t_2 \leftarrow \text{offset}(n)$ 
10     $r \leftarrow \text{NextRegister}()$ 
11    emit(loadAO,  $t_1$ ,  $t_2$ , r)
12  else if  $n = \text{num}$  then
13     $r \leftarrow \text{NextRegister}()$ 
14    emit(loadI, n,  $\_$ , r)
15  return r
  
```

## Arithmetic Operators

**Example** Invoking the routine `expr` on the AST for  $a - b \times c$  shown on the left produces the results shown on the right.



```

loadI  @a      ⇒ r1
loadAO rarp, r1 ⇒ r2
loadI  @b      ⇒ r3
loadAO rarp, r3 ⇒ r4
loadI  @c      ⇒ r5
loadAO rarp, r5 ⇒ r6
mult   r4, r6 ⇒ r7
add    r2, r7 ⇒ r8
  
```

The example assumes that  $a$ ,  $b$ , and  $c$  are not already in registers and that each resides in the current AR.

## Reducing Demand for Registers

The choice of storage locations has a direct impact on the **quality** of the generated code

- if `a` was in a global data area, we need another `loadI` and register to load it
- if `a` was already in a register, the two instructions to load it could be omitted

Code-shape decisions encoded into the treewalk code generator have an effect on demand for registers

- the naive code uses **eight** registers plus  $r_{arp}$
- a register allocator could rewrite the code as shown on the right
- the rewritten code uses **three** registers plus  $r_{arp}$

<code>loadI</code>	<code>@a</code>	$\Rightarrow$	$r_1$
<code>loadAO</code>	$r_{arp}, r_1$	$\Rightarrow$	$r_1$
<code>loadI</code>	<code>@b</code>	$\Rightarrow$	$r_2$
<code>loadAO</code>	$r_{arp}, r_2$	$\Rightarrow$	$r_2$
<code>loadI</code>	<code>@c</code>	$\Rightarrow$	$r_3$
<code>loadAO</code>	$r_{arp}, r_3$	$\Rightarrow$	$r_3$
<code>mult</code>	$r_2, r_3$	$\Rightarrow$	$r_2$
<code>add</code>	$r_1, r_2$	$\Rightarrow$	$r_2$

## Reducing Demand for Registers

A different code shape can reduce the demand for registers. Using a **right-to-left** tree walk instead of a **left-to-right** tree walk produces the code shown below on the left.

```
loadI  @c      ⇒ r1
loadAO rarp, r1 ⇒ r2
loadI  @b      ⇒ r3
loadAO rarp, r3 ⇒ r4
mult   r2, r4  ⇒ r5
loadI  @a      ⇒ r6
loadAO rarp, r6 ⇒ r7
add    r7, r5  ⇒ r2
```

```
loadI  @c      ⇒ r1
loadAO rarp, r1 ⇒ r1
loadI  @b      ⇒ r2
loadAO rarp, r2 ⇒ r2
mult   r1, r2  ⇒ r1
loadI  @a      ⇒ r2
loadAO rarp, r2 ⇒ r2
add    r2, r1  ⇒ r1
```

After register allocation, the code shown above on the right only uses **two** registers plus  $r_{arp}$ .



## Reducing Demand for Registers

Of course, right-to-left evaluation is not a general solution. To generate an evaluation order that reduces demand for registers, the compiler must choose between right and left children.

### Rule

The compiler can **minimize** register use by evaluating first, at each node, the subtree that needs the most registers.

This approach requires **two** passes over the AST

- the first pass computes the demand for registers
- the second pass emits the actual code

## Accessing Parameter Values

Formal parameters need may a different treatment

- a **call-by-value** parameter passed in the AR can be treated as if it was a local variable
- a **call-by-reference** parameter passed in the AR requires one additional indirection

**Example** For the call-by-reference parameter  $d$ , the compiler might generate the following instructions to obtain  $d$ 's value.

```
loadI    @d           $\Rightarrow$   $r_1$   
loadAO    $r_{arp}$ ,  $r_1$   $\Rightarrow$   $r_2$   
load      $r_2$           $\Rightarrow$   $r_3$ 
```

The first two instructions move the address of  $d$  into  $r_2$ , the third instruction loads its value.

## Accessing Parameter Values

Many linkage conventions pass the first few parameters in registers. So far, our code generation algorithm cannot handle a value that is **permanently** kept in a register.

The necessary extensions are easy to implement

- **Call-by-Value Parameters**

The ident case must check if the value is already in a register

- if so, it just assigns the register number to `r`
- otherwise, it uses the standard mechanisms to load the value from memory

- **Call-by-Reference Parameters**

- if the address resides in a register, simply load the value into a register
- if the address resides in the AR, must load the address before loading the value

**The Small Print** Note that the compiler cannot keep the value of a call-by-reference parameter in a register across an assignment, unless the compiler can prove that the reference is unambiguous, across all calls to the procedure.

## Function Calls in an Expression

Apart from variables, constants, and temporary values produced by other subexpressions, **function calls** also occur as operands in expressions.

To evaluate a function call, the compiler performs the following steps

- generate the calling sequence needed to invoke the function
- emit the code necessary to move the returned value to a register

**Recall** The linkage convention limits the callee's impact on the caller.

The presence of a function call may restrict the compiler's ability to **change** the evaluation order of an expression

- function may have **side effects** that modify values of variables used in the expression
- without further analysis, compiler must emit code that is correct in the **worst case**

## Other Arithmetic Operators

To handle other arithmetic operations, we can extend the treewalk model.

The basic scheme remains the same

- get the operands into registers
- perform the operation
- store the result

**Note** Operator precedence, encoded into the expression grammar, ensures the correct evaluation order.

Some operators require complex **multioperation sequences** for their implementation, e.g., exponentiation and trigonometric functions

- expand operation sequence inline
- emit call to a library routine

## Mixed-Type Expressions

Many programming languages support operations with operands of **different types**. The compiler must recognize this situation and insert the **conversion code**.

### Built-in Types

- definition of programming language specifies a formula for each conversion
- some processors provide explicit conversion operators
- others expect the compiler to generate complex, machine-dependent code

### Programmer-defined Types

- compiler has no conversion tables that define each specific case
- source language still defines the meaning of the expression
- compiler must implement this meaning and reject illegal expressions

## Assignment as an Operator

Most Algol-like languages implement assignment with the following simple rules

1. evaluate the right-hand side of the assignment to a **value**
2. evaluate the left-hand side of the assignment to a **location**
3. store the right-hand side value into the left-hand side location

Distinguishing between these modes of evaluation

- **rvalue**: result (value) of evaluation on the right-hand side of an assignment
- **lvalue**: result (address) of evaluation on the left-hand side of an assignment

In an assignment, the type of the lvalue can differ from the type of the rvalue, which may require either a compiler-inserted **conversion** or an **error message**.

**Note** The typical rule for such a conversion is to evaluate the rvalue to its natural type and then convert the result to the type of the lvalue.

## Boolean and Relational Operators

Most programming languages can operate on the results of **Boolean** and **relational operators**, both of which produce Boolean values.

To support such values, a compiler writer must...

- augment standard expression grammar with Boolean and relational operators
- define and enforce typing and inference rules for these operators
- decide how to represent Boolean values and how to compute them

Most architectures provide a **rich set** of Boolean operations, but support for relational operators **varies widely** from one architecture to another.

The compiler writer must find an evaluation strategy that matches the needs of the language to the available instruction set.



## Boolean and Relational Operators

0	$Expr \rightarrow Expr \vee AndTerm$
1	$AndTerm$
2	$AndTerm \rightarrow AndTerm \wedge RelExpr$
3	$RelExpr$
4	$RelExpr \rightarrow RelExpr < NumExpr$
5	$RelExpr \leq NumExpr$
6	$RelExpr = NumExpr$
7	$RelExpr \neq NumExpr$
8	$RelExpr \geq NumExpr$
9	$RelExpr > NumExpr$
10	$NumExpr$

11	$NumExpr \rightarrow NumExpr + Term$
12	$NumExpr - Term$
13	$Term$
14	$Term \rightarrow Term \times Value$
15	$Term \div Value$
16	$Factor$
17	$Value \rightarrow \neg Factor$
18	$Factor$
19	$Factor \rightarrow ( Expr )$
20	$num$
21	$ident$

## Representations

Traditionally, two representations have been proposed for boolean values.

### Numerical Encoding

- assign specific values to `true` and `false`
- manipulate them using the target machine's arithmetic and logical operations

### Positional Encoding

- encode the value of the expression as a position in the executable code
- use comparisons and conditional branches to evaluate the expression
- different control-flow paths represent the result of evaluation

→ Each approach works well for some examples, but not for others...

## Numerical Encoding

When the program **stores** the result of a Boolean or relational operation into a variable, the compiler must assign **numerical values** to true and false.

Typical Values that work with hardware operations such as and, or, and not

- false: zero
- true: one, word of ones, or  $\neg$ false

**Example** If b, c, and d are all in registers, the compiler might produce the following code for the expression  $b \vee c \wedge \neg d$ .

```
not  rd       $\Rightarrow$  r1
and  rc, r1  $\Rightarrow$  r2
or   rb, r2  $\Rightarrow$  r3
```

## Numerical Encoding

For a **comparison**, such as  $a < b$ , the compiler must generate code that compares  $a$  and  $b$  and assigns the appropriate value to the result.

1. the target machine provides a comparison operation that **returns a Boolean**

$$\text{cmp\_LT } r_a, r_b \Rightarrow r_1$$

2. the comparison defines a **condition code** that must be read with a branch

```

      comp      r_a, r_b  ⇒  cc1
      cbr_LT    cc1      →  L1, L2
L1: loadI     true      ⇒  r1
      jumpI                    →  L3
L2: loadI     false     ⇒  r1
      jumpI                    →  L3
L3: nop
```

## Condition Codes on Intel x86-64

Nearly all arithmetic instructions set condition codes based on their result.

**ZF** result was **z**ero

**CF** result caused **c**arry out of most significant bit

**SF** result was negative (**s**ign bit was set)

**OF** result caused (signed) **o**verflow

Based on these condition codes, standard condition suffixes *cc* are defined (*cf.* Slide 349) that modify the behaviour of three different kinds of instructions.

**jcc** jumps to the specified label if *cc* holds

**setcc** sets the given (single-byte) register to 1 or 0 depending on whether *cc* holds or not

**cmovcc** performs the specified move only if *cc* holds

The table on the right shows the **standard condition suffixes** defined by the Intel x86-64 instruction set.

**Example** The expression  $a < b$  could be implemented as follows.

```
1 cmp    %rdi, %rsi
2 setl   %al
```

**Note** Since `setl` needs a single-byte register, the example uses `%al` instead of `%rax`.

cc	Condition Tested	Meaning
e	ZF	equal to zero
ne	$\neg$ ZF	not equal to zero
s	SF	negative
ns	$\neg$ SF	not negative
g	$\neg(\text{SF} \oplus \text{OF}) \wedge \neg \text{ZF}$	greater (signed $>$ )
ge	$\neg(\text{SF} \oplus \text{OF})$	greater or equal (signed $\geq$ )
l	$\text{SF} \oplus \text{OF}$	less (signed $<$ )
le	$(\text{SF} \oplus \text{OF}) \vee \text{ZF}$	less or equal (signed $\leq$ )
a	$\neg \text{CF} \wedge \neg \text{ZF}$	above (unsigned $>$ )
ae	$\neg \text{CF}$	above or equal (unsigned $\geq$ )
b	CF	below (unsigned $<$ )
be	$\text{CF} \vee \text{ZF}$	below or equal (unsigned $\leq$ )

## Positional Encoding

Positional encoding makes sense if an expression's result is **never** stored

- result of subexpression evaluations
- expressions to determine control flow

### Short-Circuit Evaluation

In **short-circuit evaluation**, expressions are only evaluated until their final value is determined. Short-circuit evaluation relies on two Boolean identities.

$$\forall x \text{ false} \wedge x = \text{false}$$

$$\forall x \text{ true} \vee x = \text{true}$$

Some programming languages, e.g., C and Java, **require** the compiler to use short-circuit evaluation.

## Positional Encoding

**Example** Consider the following code for the expression  $a < b \vee c < d \wedge e < f$  that a naive code generator would emit.

comp $r_a, r_b \Rightarrow cc_1$	L <sub>5</sub> : loadI false $\Rightarrow r_2$
cbr_LT $cc_1 \rightarrow L_1, L_2$	jumpI $\rightarrow L_6$
L <sub>1</sub> : loadI true $\Rightarrow r_1$	L <sub>6</sub> : comp $r_e, r_f \Rightarrow cc_3$
jumpI $\rightarrow L_3$	cbr_LT $cc_3 \rightarrow L_7, L_8$
L <sub>2</sub> : loadI false $\Rightarrow r_1$	L <sub>7</sub> : loadI true $\Rightarrow r_3$
jumpI $\rightarrow L_3$	jumpI $\rightarrow L_9$
L <sub>3</sub> : comp $r_c, r_d \Rightarrow cc_2$	L <sub>8</sub> : loadI false $\Rightarrow r_3$
cbr_LT $cc_2 \rightarrow L_4, L_5$	jumpI $\rightarrow L_9$
L <sub>4</sub> : loadI true $\Rightarrow r_2$	L <sub>9</sub> : and $r_2, r_3 \Rightarrow r_4$
jumpI $\rightarrow L_6$	or $r_1, r_4 \Rightarrow r_5$

**Note** Every path takes **eleven** operations, including **three** branches and **three** jumps!



## Positional Encoding

If the compiler avoids storing intermediate results of subexpressions and uses short-circuit evaluation, it can emit **much more efficient** code.

```

      comp    ra, rb ⇒ cc1
      cbr_LT  cc1   → L3, L1
L1:  comp    rc, rd ⇒ cc2
      cbr_LT  cc2   → L2, L4
L2:  comp    re, rf ⇒ cc3
      cbr_LT  cc3   → L3, L4
L3:  loadI   true   ⇒ r5
      jumpI           → L5
L4:  loadI   false  ⇒ r5
      jumpI           → L5
L5:  nop

```

## Positional Encoding

When the code uses the result of an expression to determine control flow, positional encoding often avoids extraneous operations.

```

1 if (a < b)
2 {
3     statement1;
4 }
5 else
6 {
7     statement2;
8 }

```

```

      comp    ra, rb ⇒ cc1
      cbr_LT  cc1    → L1, L2
L1: code for statement1
      jumpI           → L3
L2: code for statement2
      jumpI           → L3
L3: nop

```

**Note** The code combines the evaluation of  $a < b$  with the selection between *statement*<sub>1</sub> and *statement*<sub>2</sub>. The result of  $a < b$  is represented as a position, either L<sub>1</sub> or L<sub>2</sub>.

## Hardware Support for Relational Operations

Specific, low-level details in the target machine's instruction set strongly influence the choice of a representation for relational values.

We will consider four schemes for supporting relational expressions

- straight condition codes
- condition codes augmented with a conditional move operation
- boolean-valued comparisons
- predicated operations

For each of these, we will examine the implementation of an `if-then-else` statement and an assignment of a Boolean value.

**Note** The first two schemes are supported by the Intel x84-64 instruction set.

<b>Source Code</b>	<pre> if (x &lt; y)   then a ← c + d   else a ← e + f </pre>	
<b>ILOC Code</b>	<pre>       comp   r<sub>x</sub>, r<sub>y</sub> ⇒ cc<sub>1</sub>       cbr_LT cc<sub>1</sub>  → L<sub>1</sub>, L<sub>2</sub>  L<sub>1</sub>: add     r<sub>c</sub>, r<sub>d</sub> ⇒ r<sub>a</sub>       jumpI           → L<sub>out</sub>  L<sub>2</sub>: add     r<sub>e</sub>, r<sub>f</sub> ⇒ r<sub>a</sub>       jumpI           → L<sub>out</sub>  L<sub>out</sub>: nop        Straight Condition Codes </pre>	<pre>       cmp_LT  r<sub>x</sub>, r<sub>y</sub> ⇒ r<sub>1</sub>       cbr     r<sub>1</sub>    → L<sub>1</sub>, L<sub>2</sub>  L<sub>1</sub>: add     r<sub>c</sub>, r<sub>d</sub> ⇒ r<sub>a</sub>       jumpI           → L<sub>out</sub>  L<sub>2</sub>: add     r<sub>e</sub>, r<sub>f</sub> ⇒ r<sub>a</sub>       jumpI           → L<sub>out</sub>  L<sub>out</sub>: nop        Boolean Compare </pre>
	<pre> comp   r<sub>x</sub>, r<sub>y</sub>      ⇒ cc<sub>1</sub> add    r<sub>c</sub>, r<sub>d</sub>      ⇒ r<sub>1</sub> add    r<sub>e</sub>, r<sub>f</sub>      ⇒ r<sub>2</sub> i2i_LT cc<sub>1</sub>, r<sub>1</sub>, r<sub>2</sub> ⇒ r<sub>a</sub>        Conditional Move </pre>	<pre>       cmp_LT  r<sub>x</sub>, r<sub>y</sub> ⇒ r<sub>1</sub>       not     r<sub>1</sub>    ⇒ r<sub>2</sub> (r<sub>1</sub>)? add    r<sub>c</sub>, r<sub>d</sub> ⇒ r<sub>a</sub> (r<sub>2</sub>)? add    r<sub>e</sub>, r<sub>f</sub> ⇒ r<sub>a</sub>        Predicated Execution </pre>

Source Code	$x \leftarrow a < b \wedge c < d$	
<b>ILOC Code</b>	<pre>       comp  r<sub>a</sub>, r<sub>b</sub>  ⇒ cc<sub>1</sub>       cbr_LT cc<sub>1</sub>   → L<sub>1</sub>, L<sub>2</sub>  L<sub>1</sub>:  comp  r<sub>c</sub>, r<sub>d</sub>  ⇒ cc<sub>2</sub>       cbr_LT cc<sub>2</sub>   → L<sub>3</sub>, L<sub>2</sub>  L<sub>2</sub>:  loadI false ⇒ r<sub>x</sub>       jumpI      → L<sub>out</sub>  L<sub>3</sub>:  loadI true  ⇒ r<sub>x</sub>       jumpI      → L<sub>out</sub>  L<sub>out</sub>: nop        Straight Condition Codes </pre>	<pre> comp  r<sub>a</sub>, r<sub>b</sub>      ⇒ cc<sub>1</sub> i2i_LT cc<sub>1</sub>, r<sub>T</sub>, r<sub>F</sub> ⇒ r<sub>1</sub> comp  r<sub>c</sub>, r<sub>d</sub>      ⇒ cc<sub>2</sub> i2i_LT cc<sub>2</sub>, r<sub>T</sub>, r<sub>F</sub> ⇒ r<sub>2</sub> and   r<sub>1</sub>, r<sub>2</sub>      ⇒ r<sub>x</sub>        Conditional Move </pre>
		<pre> cmp_LT r<sub>a</sub>, r<sub>b</sub>      ⇒ r<sub>1</sub> cmp_LT r<sub>c</sub>, r<sub>d</sub>      ⇒ r<sub>2</sub> and    r<sub>1</sub>, r<sub>2</sub>      ⇒ r<sub>x</sub>        Boolean Compare </pre>
		<pre> cmp_LT r<sub>a</sub>, r<sub>b</sub>      ⇒ r<sub>1</sub> cmp_LT r<sub>c</sub>, r<sub>d</sub>      ⇒ r<sub>2</sub> and    r<sub>1</sub>, r<sub>2</sub>      ⇒ r<sub>x</sub>        Predicated Execution </pre>

## Hardware Support for Relational Operations

### Condition Codes

- code has at least one conditional branch per relational operator
- comparison operation may be omitted if condition codes are set by default

### Conditional Move

- leads to faster code by avoiding branches
- safe as long as neither operation can raise an exception

### Boolean Compare

- works without a branch and without converting comparison results to Boolean values
- a weakness of this model is that it requires explicit comparisons

### Predicated Execution

- code is simple and concise
- predication can lead to the same code as the boolean-comparison scheme

## Relational Operations on Intel x86-64

The Intel x86-64 instruction set supports both the **straight condition codes** and **conditional move** implementation scheme.

Straight condition codes are typically used to implement control-flow constructs, whereas conditional moves are used to evaluate and assign expressions.

A compiler for Intel x86-64 might implement the two examples from before as follows.

```
1      cmpq    %rdi, %rsi
2      jl     _L1
3      leaq    (%rdx, %rcx), %rax
4      jmp     _L2
5 _L1:  leaq    (%r8, %r9), %rax
6 _L2:  nop
```

```
1      cmpq    %rdi, %rsi
2      setl    %al
3      cmpq    %rdx, %rcx
4      setl    %bl
5      and     %al, %bl
6      movzbq  %bl, %rax
```